

École polytechnique de Louvain

Observing the detailed behaviour of large distributed applications in real time using Δ QSD

The Δ Q Oscilloscope for Erlang applications

Author: **Francesco NIERI**

Supervisor: **Peter VAN ROY**

Readers: **Neil DAVIES, Tom BARBETTE, Peer STRITZINGER**

Academic year 2024–2025

Master [120] in Computer Science

Abstract

It is difficult to study the detailed behaviour of large distributed systems while they are running. What happens when there is an overload? Modern software development practices successfully fail to adequately consider essential quality requirements or even to consider properly whether a system can actually meet its intended outcomes. Current observability tools do not meet observability requirements when it comes to detecting problems early enough in running systems.

To tackle these problems, PNSol Ltd. developed the Δ QSD paradigm, a novel metrics-based and quality-centric paradigm that uses formalised outcome diagrams to explore the performance consequences of design decisions, as a performance blueprint of the system. PNSol Ltd. analyses the behaviour of existing systems using Δ QSD, but the technology only works a posteriori, there is no way yet to analyse a system's behaviour in real time.

To advance the usage of Δ QSD in distributed projects, this project aims to develop a prototype of the Δ Q oscilloscope, a dashboard to observe a running Erlang system in real time. The oscilloscope works by communicating to a wrapper which attaches probes to an Erlang system, the "system under test", which translates OpenTelemetry spans to outcome instances and sends them to the oscilloscope.

In order to use the Δ Q Oscilloscope, the first step is to determine the system's outcome diagram. The system's outcome diagram is a directed graph that shows the causal relationships between the system's internal operations (called outcomes) and its overall inputs and outputs. Each outcome corresponds to a system component. It has a start event and stop event, when the component is called and when it returns.

For each outcome of interest, an observation point (probe) is attached to measure the delay of that outcome. As many probes are put into the system as are needed to observe the desired system behaviours. The probes are implemented on top of OpenTelemetry tracing ability. Each probe sends a time series of data to the Δ Q Oscilloscope, which performs statistical computations on all the time series and displays the results in real time.

The oscilloscope provides a system of triggers, which, like a true oscilloscope, captures rare events in a distributed system. The oscilloscope can provide snapshots of the running system under test which provide a still of the system, as if it was frozen in time.

AI disclaimer

AI was used to aid the writing of the code. The **written master thesis** was written entirely without the aid of AI.

Contents

1	Introduction	1
1.1	Context	1
1.2	Objective	1
1.3	Previous work	2
1.4	Contributions	2
1.5	Roadmap	3
2	Background	4
2.1	An overview of Δ QSD	4
2.1.1	Outcome	4
2.1.2	Quality attenuation (Δ Q)	5
2.1.3	Failure semantics	6
2.1.4	Partial ordering	6
2.1.5	Timeliness	6
2.1.6	QTA, required Δ Q	6
2.1.7	Outcome diagram	7
2.1.8	Independence hypothesis	9
2.2	Observability	10
2.2.1	erlang:trace	11
2.2.2	OpenTelemetry	11
2.3	Current observability problems	13
2.3.1	Handling of long span	14
3	Design	15
3.1	Probes	15
3.2	Extending the notion of failure	16
3.3	Time series - Oscilloscope outcome instances	16
3.4	Erlang system	17
3.4.1	System under test	17
3.4.2	Wrapper/Adapter	17
3.4.3	Inserting probes in Erlang - From spans to outcome instances	18
3.5	Oscilloscope: C++ system	18
3.5.1	Server	18
3.5.2	Oscilloscope	18
3.5.3	Inserting probes in the oscilloscope	18

3.6	Oscilloscope system design diagram	19
3.7	Triggers	20
3.7.1	Snapshot	20
3.8	Polling window	20
4	Oscilloscope: User level concepts	22
4.1	Δ QSD implementation	22
4.1.1	Internal representation of a Δ Q	22
4.1.2	dMax	23
4.1.3	QTA	24
4.1.4	Confidence bounds	24
4.2	Δ Q display	24
4.3	Outcome diagram	25
4.3.1	Probes	25
4.3.2	Operators, outcome expressions	25
4.3.3	Limitations	26
4.4	Dashboard	26
4.4.1	Sidebar	26
4.4.2	Plots window	28
4.4.3	Stub controls	28
4.5	Triggers	28
4.5.1	Load	28
4.5.2	QTA	28
5	Oscilloscope: implementation	29
5.1	Δ QSD implementation	29
5.1.1	dMax	30
5.1.2	Operations	30
5.1.3	Confidence bounds	31
5.1.4	Rebinning	31
5.2	Wrapper	32
5.2.1	API	32
5.2.2	Handling outcome instances	34
5.2.3	TCP connection	35
5.3	Parser	35
5.3.1	ANTLR	36
5.3.2	Grammar	36
5.4	Oscilloscope GUI	37
6	Application on synthetic programs	39
6.1	Sequential composition with M/M/1/K queues	39
6.1.1	System composition	39
6.1.2	Determining parameters dynamically	40
6.1.3	Triggers - QTA	42
6.2	First to finish application	44
6.2.1	Using the wrong operator	44
6.2.2	Introducing a slower component	45

7	Performance study	46
7.1	Convolution performance	46
7.2	Stub performance	47
7.3	GUI plotting performance	47
8	Conclusions and future work	49
8.1	Future improvements	49
8.1.1	Oscilloscope improvement	50
8.1.2	Wrapper improvements	50
8.1.3	Real applications	51
8.1.4	Licensing limitations	51

Chapter 1

Introduction

1.1 Context

Δ QSD is an industrial-strength approach for large-scale system design that can predict performance and feasibility early on in the design process. Developed over 30 years by a small group of people around Predictable Network Solutions Ltd, the paradigm has been applied in various industrial-scale problems with huge success and large savings in costs. [1] Modern software development practices successfully fail to adequately consider essential quality requirements or even to consider properly whether a system can actually meet its intended outcomes, particularly when deployed at scale, the Δ QSD paradigm addresses this problem!

Δ QSD has important properties which make its application to distributed projects interesting, it supports:

- A compositional approach that considers performance and failure as first-class citizens.
- Stochastic approach to capture uncertainty throughout the design approach.
- Performance and feasibility can be predicted at high system load for partially defined systems.

While the paradigm has been successfully applied in **a posteriori** analysis, there is no way yet to analyse a distributed system which is running in real time! This is where the Δ Q oscilloscope comes in.

1.2 Objective

This project will develop a practical tool, the **Δ Q scilloscope**, for the Erlang developer community.

The Erlang language and Erlang/OTP platform are widely used to develop distributed applications that must perform reliably under high load. The tool will provide useful

information for these applications both for understanding their behaviour, for diagnosing performance issues, and for optimizing performance over their lifetime.

The ΔQ Oscilloscope will perform statistical computations to show real time graphs about the performance of system components. With the oscilloscope prototype we will present in this paper, we are aiming to show that the ΔQSD paradigm is not only a theoretical paradigm, but it can be employed in a tool to diagnose large distributed systems.

The oscilloscope targets large distributed applications handling many independent tasks where performance and reliability are important.

1.3 Previous work

The ΔQSD paradigm has been formalised across different papers [2] [3] and was brought to the attention of engineers via tutorials [1] and to students at Université Catholique de Louvain [4].

A Jupyter notebook workbench has been made available on GitHub [5], it shows real time ΔQ graphs for typical outcome diagrams but is not adequate to be scaled to real time systems, it is meant as an interactive tool to show how the ΔQSD paradigm can be applied to real life examples.

Observability tools such as Erlang tracing [6] and OpenTelemetry [7] lack the notions of failure as defined in ΔQSD , which allows detecting performance problems early on, we base our program on OpenTelemetry to incorporate already existing notions of causality and observability to augment their capabilities and make them suitable to work with the ΔQSD paradigm.

1.4 Contributions

There are a few contributions that make the master thesis and thus, the oscilloscope, possible:

- A graphical interface to display ΔQ plots for outcomes.
- An Erlang OpenTelemetry wrapper to give OpenTelemetry spans a notion of failure and to communicate with the oscilloscope.
- An implementation of a syntax, derived from the original algebraic syntax to create outcome diagrams.
- The implementation of ΔQSD concepts from theory to practice, allowing probes' $\Delta Q(s)$ to be displayed and analysed on the oscilloscope.
- An efficient convolution algorithm based on the FFTW3 library.
- A system of triggers to catch rare events when system behaviour fails to meet quality requirements, giving a snapshot of the system, giving the user insights about their system's behaviour.

- Synthetic applications to test the effectiveness of Δ QSD on diagnosing systems and their feasibility.

These contributions can show that the Δ QSD has its practical applications and is not limited to a theoretical view of system design.

1.5 Roadmap

The following thesis will give the reader everything that is needed to use the oscilloscope and exploit it to its full potential.

We divided the thesis in multiple chapters, below is the roadmap of the content:

- The background chapter gives the reader an extensive background into the theoretical foundations of Δ QSD, which are the basis of the oscilloscope and are fundamental to understand how to correctly use and analyse the output given by the oscilloscope. Secondly, an introduction to OpenTelemetry, the library we base our Erlang adapter on, and the problems that are present in the observability library.
- The design chapter initially extends the concepts of Δ QSD, introducing novel aspects which are key to understand how to properly instrument the different part of the systems, as they are the basis upon which we build the oscilloscope. We then delve how the parts of the system interact together and how to correctly apply the concepts we just introduced in the oscilloscope and the Erlang system.
- We then present the oscilloscope in two different chapters, first providing "user level concepts" of how Δ QSD is used in the oscilloscope and what the user should expect graphically from the oscilloscope. Secondly, a more low level explanation, which goes into more technical details of the parts that compose the oscilloscope and the Δ QSD concepts explained in the previous chapter.
- We then provide synthetic applications which have been tested with the oscilloscope that demonstrate the usefulness of the oscilloscope in a distributed setting. We also perform evaluations of the performance of the different parts we have developed to understand the overhead that are present.

We end by providing future possibilities which can be explored, and concepts which we believe ought to be implemented in observabilities tools. In the appendix, we provide a user manual to help users use the oscilloscope, along with C++ and Erlang source code of the oscilloscope and the wrapper.

Chapter 2

Background

This chapter aims to provide firstly a complete background of the concepts key to understanding the Δ QSD, giving a good basis to grasp the concepts of **quality attenuation** (Δ Q) and outcome diagrams (SD).

Secondly, we provide a comprehensive background into the observability solutions that have been explored for the oscilloscope, delving deeper into OpenTelemetry and its macros.

We finish by explaining the current limitations of OpenTelemetry and explaining where our oscilloscope comes in.

2.1 An overview of Δ QSD

Δ QSD is a metrics-based, quality-centric paradigm that uses formalised outcome diagrams to explore the performance consequences of design decisions. [2]

Key concepts of Δ QSD are **quality attenuation** (Δ Q) and **outcome diagram**

Outcome diagrams capture observational properties of the system. The Δ QSD paradigm derives bounds on performance expressed as probability distribution, encompassing all possible executions of the system.

The following sections are a summary of multiple articles and presentation formalizing the paradigm. [2] [1] [3] [8]

2.1.1 Outcome

An outcome O is a specific system behaviour that can be observed to start at some point in time and **may** be observed to complete at some later time. [8] Formally, what the system obtains by performing one of its tasks. One task corresponds to one outcome and viceversa. When an outcome is performed, it means that the task of an outcome is performed.

Assume a component C which receives a message m_{in} and outputs a message m_{out} after a delay d . Over multiple executions, we will have observed multiple delays which can

be represented as a cumulative definition where p percent of delays have delay $\leq d$, the ΔQ . [3]

Observables Each outcome has two starting sets of events: the starting sets and the ending sets. Such sets are called the *observables*. Once an event from the starting set occurs, there is no guarantee that a corresponding event in the terminating set will occur within the duration limit (required time to complete). An observable is *done* when it occurs during the time limit.

Outcome instance Given a starting event e_{in} and an end event e_{out} , an *outcome instance* is what the system gains within (e_{in}, e_{out}) .

Representation Outcome are represented as circles, with the starting and terminating set of events being represented by boxes.




Figure 2.1: The outcome (circle) and the starting set (left) and terminating set (right) of events.

2.1.2 Quality attenuation (ΔQ)

ΔQ is a cumulative distribution function that defines both *latency* and *failure probability* between a start and end event [1]

In an ideal system, an outcome would deliver a desired behaviour without error, failure, delay, but this is not the case. The quality of an outcome response "attenuated to the relative ideal" (the cumulative distribution function) is called "quality attenuation" (ΔQ). Since it can depend on many factors (geographical, physical ...), ΔQ is modeled as a random variable.

As ΔQ captures deviation from ideal behavior, and incorporates delay, which is a continuous random variable, and failures/timeouts, which are discrete variables, it can be described mathematically as an *Improper Random Variable*, where the probability of a delay < 1 . Combining latency and failure together makes it easy to examine the tradeoffs between them.

$\Delta Q(x)$ is the probability that an outcome O occurs in time $t \leq x$. The *intangible mass* $1 - \lim_{x \rightarrow \infty} \Delta Q(x)$ of a ΔQ will encode the probability of failure/timeout/exception occurring.




Figure 2.2: Intangible mass (red) of a ΔQ , the observable had a failure rate of about 5%

2.1.3 Failure semantics

In the CDF representation of a ΔQ , there is an f percent probability that the delay is infinite, this is what failure models. Concretely, it means that an input message m_{in} **has no output message** m_{out} .

Combining delay and failure in a single quantity is what makes ΔQSD a great choice to explore feasibility in system design.

2.1.4 Partial ordering

A CDF of a ΔQ is *less than* the other if its CDF is everywhere to the left and above the other. Mathematically, it is a partial order.

If two ΔQ s intersect, they are not ordered.

2.1.5 Timeliness

Timeliness is defined as a relation between an observed ΔQ_{obs} and a required ΔQ_{req} . Timeliness is delivering results within required time bounds (sufficiently often).

A system *satisfies timeliness* if $\Delta Q_{obs} \leq \Delta Q_{req}$.

2.1.6 QTA, required ΔQ

The *Quantitative Timeliness Agreement* (QTA) maps objective measurements to the subjective perception of application performance [8]. It specifies what the base system does and its limits.

Slack When ΔQ is strictly less than the requirement, we say there is performance *slack*.

Hazard When ΔQ is strictly greater than the requirement, there is performance hazard.

QTA example : Imagine a system where 25% of the executions should take < 15 ms, 50% < 25 ms and 75% < 35 ms, all queries have a maximum delay of 50ms and 5% of executions can timeout, the QTA can be represented as a step function.

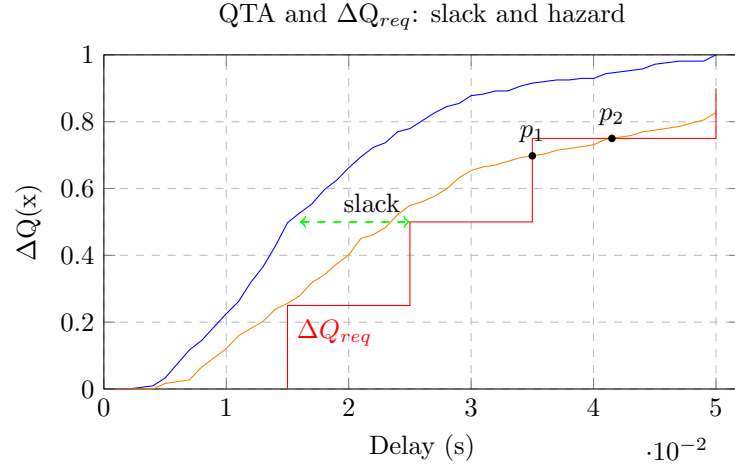


Figure 2.3: The system in blue is showing slack and satisfies the requirement, the system in orange is showing signs that it cannot handle the stress, it is not respecting the ΔQ_{req} .

2.1.7 Outcome diagram

An outcome diagram captures the causal relationships between the outcomes, each outcome diagram can be presented algebraically with an outcome expression. It allows computing the ΔQ for the whole system. The outcome diagram should capture the essential observational properties of a system. There are four different ways to represent the relationships between outcomes.

Sequential composition

If we assume two outcomes O_A , O_B where end event of O_A is the start event of O_B , the two outcomes can be sequentially composed. The ΔQ of O_{AB} is given by the convolution of the PDFs of O_A and O_B .




Figure 2.4: Sequential composition of O_A and O_B .

Where convolution (\circledast) between two PDF is :

$$PDF_{AB}(t) = \int_0^t PDF_A(\delta) \cdot PDF_B(t - \delta) d\delta \quad (2.1)$$

and thus ΔQ_{AB} :

$$\Delta Q_{AB} = \Delta Q_A \circledast \Delta Q_B \quad (2.2)$$

First to finish

If we assume two independent outcomes O_A, O_B with the same start event, first-to-finish occurs when at least one end event occurs, it can be calculated as:

$$\begin{aligned} \Delta Q_{FTF(A,B)} &= Pr[d_A > t \wedge d_B > t] \\ &= Pr[d_A > t] \cdot Pr[d_B > t] = (1 - \Delta Q_A) \cdot (1 - \Delta Q_B) \\ \Delta Q_{FTF(A,B)} &= \Delta Q_A + \Delta Q_B - \Delta Q_A \cdot \Delta Q_B \end{aligned} \quad (2.3)$$




Figure 2.5: Left: Two ΔQ s, A and B. Right: The result of the first to finish operator applied on A and B

All to finish

If we assume two independent outcomes O_A, O_B with the same start event, all-to-finish occurs when both end events occur, it can be calculated as:

$$\begin{aligned} \Delta Q_{ATF(A,B)} &= Pr[d_A \leq t \wedge d_B \leq t] \\ &= Pr[d_A \leq t] \cdot Pr[d_B \leq t] = \Delta Q_A \cdot \Delta Q_B \\ \Delta Q_{ATF(A,B)} &= \Delta Q_A \cdot \Delta Q_B \end{aligned} \quad (2.4)$$




Figure 2.6: Left: Two ΔQ s, A and B. Right: The result of the all to finish operator applied on A and B

Probabilistic choice

If we assume two possible outcomes O_A and O_B and exactly one outcome is chosen during each occurrence of a start event and:

- O_A happens with probability $\frac{p}{p+q}$
- O_B happens with probability $\frac{q}{p+q}$

$$\Delta Q_{PC}(A, B) = \frac{p}{p+q} \Delta Q_A + \frac{q}{p+q} \Delta Q_B \quad (2.5)$$





Figure 2.7: The possible operators in an outcome diagram: Probabilistic choice, first-to-finish, all-to-finish

First-to-finish, All-to-finish and probabilistic-choice are calculated on the CDF of the ΔQ .

These outcome expressions can be assembled together to create an outcome diagram, later on, we will see how one can put translate the graphical representation to outcome diagrams which can be used in the ΔQ oscilloscope.

2.1.8 Independence hypothesis

Assume two sequentially composed outcomes o_1, o_2 running on the same processor. A probe p observing the execution from the start event of o_1 to the end event of o_2 .



At low load, the two components behavior will be independent, the system will behave linearly, the observed delay of the probe p will be equal to the convolution of o_1 , o_2 ($o_1 \otimes o_2$).

When load increases, the two components will start to show dependent behaviour due to the processor utilisation increasing, the observed ΔQ will then deviate from what is calculated.




Figure 2.8: When the components are independent, what is observed (blue) and calculated (red) can be superposed, whilst when o_1 and o_2 show initial signs of dependency, what is observed (green) can be seen deviating from the calculated ΔQ .

When the system is far from being overloaded, the effect is noticeable thanks to ΔQSD even if the system is far from being overloaded. As the cliff edge of overload is approached, the nonlinearity will increase.

2.2 Observability

Observability refers to the ability to understand the internal state by examining its output, in the context of a distributed system, being able to understand the internal state of the system examining its telemetry data. [9]

In the case of the Erlang programming language, we explain below two tools that can be used to observe an Erlang program.

2.2.1 erlang:trace

The Erlang programming language gives the users different ways to observe the behaviour of the code, one of those is the function `erlang:trace/3`. The erlang run-time system exposes several trace points that can be observed, observing the trace points allows users to be notified when they are triggered [6]. One can observe function calls, messages being sent and received, process being spawned, garbage collecting

```
-spec trace(PidPortSpec, How, FlagList) -> integer()
when
    PidPortSpec ::
        pid() |
        port() |
        all | processes | ports | existing | existing_processes |
        ↪ existing_ports | new |
        new_processes | new_ports,
    How :: boolean(),
    FlagList :: [trace_flag()].
```

Figure 2.9: `erlang:trace/3` specification

Nevertheless, Erlang Tracing, according to our use case, has a major flaw: no notion of causality. If two messages a, b are sent and then received in disorder, the tracer has no default way of knowing which is which, this is a missing feature that is crucial for observing a program functioning and being able to connect an application to our oscilloscope. This is where the OpenTelemetry standard comes in.

2.2.2 OpenTelemetry

OpenTelemetry is an open-source, vendor agnostic observability framework and toolkit designed to generate, export and collect telemetry data, in particular traces, metrics and logs. [9]

OpenTelemetry is available for a plethora of languages, including Erlang, although, as of writing this, only traces are available in Erlang.

OpenTelemetry provides a standard protocol, a single set of API and conventions and lets you own the generated data, allowing to switch between observability backends freely.

The Erlang Ecosystem Foundation has a working group focused on evolving the tools related to observability.

Traces

Traces are why we are basing our program on top of OpenTelemetry, traces follow the whole "path" of a request in an application, traces are comprised of one or more spans.

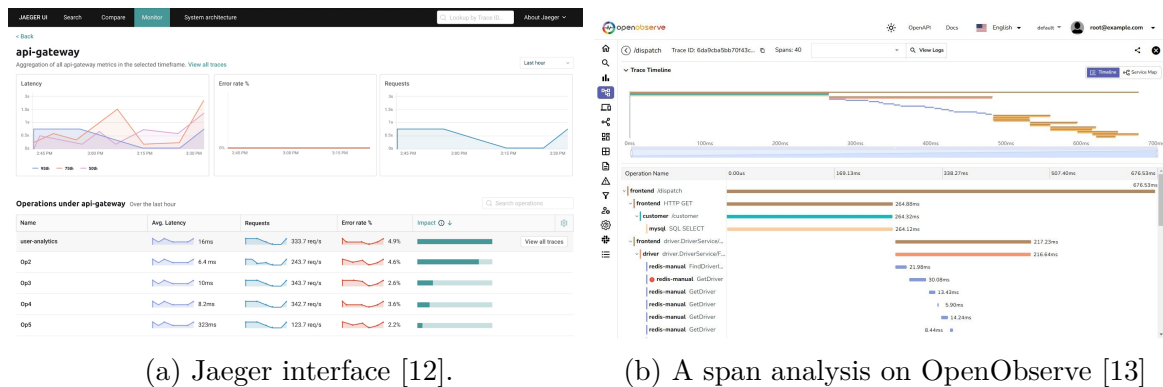
Span A span is a unit of work or operation. Spans can be correlated to each other and can be assembled into a trace. The notion of spans and traces allows us to follow the execution of a "request" and carry a context, allowing us to get the causal links of messages. [10]

```
{
  "name": "oscilloscope-span",
  "context": {
    "trace_id": "5b8aa5a2d2c872e8321cf37308d69df2",
    "span_id": "5fb397be34d26b51"
  },
  "parent_id": "0515505510cb55c13",
  "start_time": "2022-04-29T18:52:58.114304Z",
  "end_time": "2022-04-29T22:52:58.114561Z",
  "attributes": {
    "http.route": "some_route"
  }
},
}
```

Figure 2.10: Example of span with a parent, indicating a causal link between parent and children span [10]

Exporters

OpenTelemetry gives the possibility to export traces to backends such as Jaeger or Zipkin. A user can monitor their workflows, analyse dependencies, troubleshoot their programs by observing the flow of the requests in such backends. [11]



Macros

OpenTelemetry provides macros to start, end and interact with spans in Erlang, the following code excerpts are taken from the instrumentation wiki. [14]

?with_span `?with_span` creates active spans. An active span is the span that is currently set in the execution context and is considered the "current" span for the ongoing operation or thread. [15]

```

parent_function() ->
    ?with_span(parent, #{}, fun child_function/0).

child_function() ->
    %% this is the same process, so the span parent set as the active
    %% span in the with_span call above will be the active span in
    ↪ this function
    ?with_span(child, #{},
        fun() ->
            %% do work here. when this function returns, child
            ↪ will complete.
        end).

```

?start_span ?start_span creates a span which isn't connected to a particular process, it does not set the span as the current active span.

```

SpanCtx = ?start_span(child),
Ctx = otel_ctx:get_current(),
proc_lib:spawn_link(fun() ->
    otel_ctx:attach(Ctx),
    ?set_current_span(SpanCtx),
    %% do work here
    ?end_span(SpanCtx)
end),

```

?end_span ?end_span ends a span started with ?start_span

2.3 Current observability problems

A legitimate question would ask why one would need an additional tool to observe their system, observability tools are already plenty and provide useful insights into an application's behaviour. While they may seem adequate to provide a global oversight of applications, they fail to diagnose real time problems like overload, dependent behaviour early enough and in a quick manner.

The problem we are trying to tackle can be described by the following situation: Imagine an Erlang application instrumented with OpenTelemetry, suddenly, the application starts slowing down, and the execution of a function takes 10 seconds instead of the usual 1 second. Between its start and its end, the user instrumenting the application sees nothing in his dashboard.

This is a big problem! One would like to know right away if something is wrong with their application. This is where the Δ QSD paradigm and the Δ Q oscilloscope come in handy.

By extending Δ QSD notion of early failure one can now right away when the system is overloaded and showing problems, one can know right away when something is wrong,

as soon as the maximum delay is hit, this avoids waiting, in this case, 10 seconds to know that something is wrong with your application.




Figure 2.12: Execution of a long span in OpenTelemetry, the user will only be notified after 10 seconds that the function has ended (and taken too long).




Figure 2.13: Execution of a long span in OpenTelemetry, the *dMax* deadline allows knowing that the span has taken too long.

2.3.1 Handling of long span

OpenTelemetry presents a bigger problem, what happens when there are long-running spans? Worse, what happens when spans are not actually terminated?

OpenTelemetry limits the length of its spans, moreover, those who are not terminated are lost and not exported. Why? All spans must be ended!

If the span is the parent/root span, its effect could trickle down to child spans. We can quickly see how this becomes problematic, all the information about an execution of your program ...lost. Moreover, a span could not be terminated for trivial reasons: refreshing a tab, network failures, crashes ... [16]. There are a few hacks that can be implemented, having shorter spans, carrying data in child spans, saving spans in a log to track spans which were not ended to manually set an end time; why the need to circumvent limitations when observing a system?

These problems are a big limitation of OpenTelemetry, we believe that the wrapper we provide can be a great start to improve observability requirements in OpenTelemetry.

Chapter 3

Design

This chapter aims to first extend the concepts of ΔQSD , giving more insights into how the systems need to be instrumented to correctly work together, and how the different parts need to be integrated to interact together.

- We first provide concepts of probes, we extend the ΔQSD notion of failure and describe how time series will work in our oscilloscope, this part is crucial to understand how the different parts of the system work together.
- We then split the design of the oscilloscope in two. First explaining the Erlang side, where the system to be tested is. Secondly, we explain the C++ side. Both chapters explain how probes can be inserted and made to work together. We conclude the sections by showing the system diagram of the different parts.
- Lastly, we provide high level concepts of the key elements of the oscilloscope.

3.1 Probes

To observe a system, the outcomes, the result of causal links and outcome expressions, we must put probes in it.

For each outcome of interest, a probe (observation point) is attached to measure the delay of the outcome, like one would in a true oscilloscope.

Consider the figure below, a probe is attached at every component to measure their ΔQ s (c_2, c_3),

Another probe (p_1) is inserted at the beginning and end of the system to measure the global execution delay.

Thanks to this probe, the user can observe the ΔQ "*observed at p_1* ", which is the ΔQ which was calculated from the data received by inserting probe p_1 . The ΔQ "*calculated at p_1* " is the resulting ΔQ from the convolution of the observed ΔQ s at c_2 and c_3 .

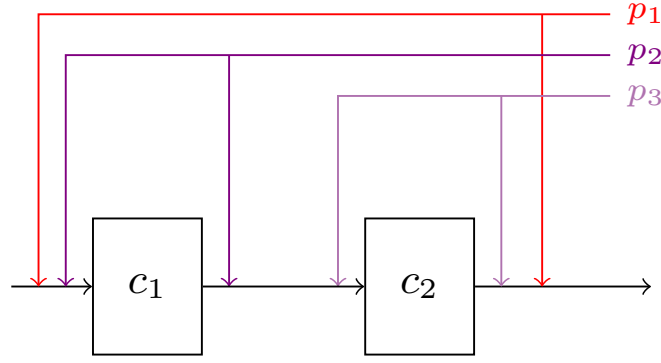


Figure 3.1: Probes inserted in a component diagram.

3.2 Extending the notion of failure

Whilst previously we defined failure as *"an input message m_{in} that has no output message m_{out} "*, we extend this definition. If you recall the previous section 2.3, we introduced the notion of a maximum delay.

We extend the notion of failure to the following definition:

"An input message m_{in} that has no output message m_{out} after x seconds"

Where x is the $dMax$ defined by the user. We can leverage this new definition to observe ΔQs in real time.

3.3 Time series - Oscilloscope outcome instances

Consider a probe p with two distinct sets of events, the starting set of events s and ending set of event e . The outcome instance of a message $m_s \rightarrow m_e$:

- The probe's p name
- The start time t_s
- The end time t_e
- Its status
- Its elapsed time of execution

The instance has three possible statuses: **success**, **timeout**, **failure**, it can thus be broken down in the representations, based on its status:

- (t_s, t_e) : This representation indicates that the execution was successful ($t < dMax$).
- (t_s, \mathcal{T}) : This representation indicates that the execution has timed out ($t > dMax$). The end time and elapsed time is equal to $t_s + \text{timeout}$
- (t_s, \mathcal{F}) : This representation indicates the execution has failed given a user defined requirement (i.e. a dropped message given buffer overload in a queue system). It

must not be confused with a program failure (crash), if a program crashes during the execution of event e , it will time out since the wrapper will not receive an end message.

The **time series** of a probe is the sequence of n outcome instances and can then be easily modeled by ΔQ .

What can be considered a failed execution? Imagine a queue with a buffer: the buffer queue being full and dropping incoming messages can be modeled as a failure.

More generally, the choice of what is considered a failed execution is left up to the user who is handling the spans and is program-dependent. Exceptions or errors can be kinds of failure.

On another note, the way of handling errored spans in OpenTelemetry can differ from user to user, so the wrapper will not handle ending and setting statuses for "failed" spans.

3.4 Erlang system

3.4.1 System under test

The system under test (**S.U.T**) is the Erlang system the engineer wishes to observe, it ideally is a system which already is instrumented with OpenTelemetry. The ideal system where ΔQSD is more useful is a system that executes many independent instances of the same action.

3.4.2 Wrapper/Adapter

The adapter is the `dqsd_otel` Erlang application, a wrapper that starts and ends OpenTelemetry spans and translates them to outcome instances which are useful for the oscilloscope. This can be done thanks to probes being attached to the system under test, like an oscilloscope would! The outcome instances end normally like OpenTelemetry spans or, additionally, can timeout, given a custom timeout, and fail, according to user's definition of failure.

Handling of OpenTelemetry spans which goes beyond starting and ending them is delegated to the user, who may wish to do further operations with their spans. The wrapper is called from the system under test and communicates outcome instances data to the oscilloscope via TCP.

The wrapper can receive messages from the oscilloscope, the messages are about updating observable's *dMax* or starting and stopping the sending of data to the oscilloscope.

3.4.3 Inserting probes in Erlang - From spans to outcome instances

OpenTelemetry spans are useful to carry context, attributes and baggage in a program. The plethora of attributes they have is nevertheless too much for the oscilloscope.

To get the equivalent of spans for the oscilloscope, the wrapper needs to be called at the starting events of a probe to start an instance of a probe, and at the ending events to end the outcome instance and send the data to the oscilloscope. The name given with "start_span" is the name of the probe.

```
% Start the outcome instance of worker_2
{WorkerCtx, WorkerPid} = otel_wrapper:start_span(<<"worker_2">>),

% Do work here ...

%End the outcome instance of worker_2
otel_wrapper:end_span(WorkerCtx, WorkerPid),
```

3.5 Oscilloscope: C++ system

3.5.1 Server

The server is responsible for receiving the messages containing the outcome instances from the wrapper. The server forwards the instances to the oscilloscope.

3.5.2 Oscilloscope

The oscilloscope receives the instances corresponding to probes from the server and adds them to the time series of the probes whose instance is being received. The oscilloscope has a graphical interface which allows the user to create an outcome diagram of the system under test, display real time graphs which show detail about the execution of the system and allow the user to set custom timeouts for probes. It can also display snapshots of the system as if it was frozen in time

3.5.3 Inserting probes in the oscilloscope

Probes are automatically inserted in the oscilloscope when creating an outcome diagram. They are inserted on the outcomes observables, outcome expressions and to the causal result of outcome expressions, we will see later on how they can be defined and how an outcome diagram can be created.

In the system below, which is equal to the one defined above, probes are automatically attached to outcomes o_1, o_2 . The user who wants to observe the result of the sequential composition can insert probes at the start and end of the routine.




Figure 3.2: Probes inserted in the outcome diagram of the previous component diagram.
Figure 3.1

As for operators (outcome expression), probes are automatically attached to the components inside them and to the start event and end events of the operators.




Figure 3.3: Probes inserted into an operator.

The **observed** ΔQ for the first-to-finish operator is the ΔQ from the instances (**start**, **end**). The **calculated** ΔQ is the ΔQ which is the result of the first-to-finish operator being applied on o_1, o_2

3.6 Oscilloscope system design diagram

Now that we have an idea of how all the parts behave together and what they do, we can draw a complete diagram of the interactions of the different components.

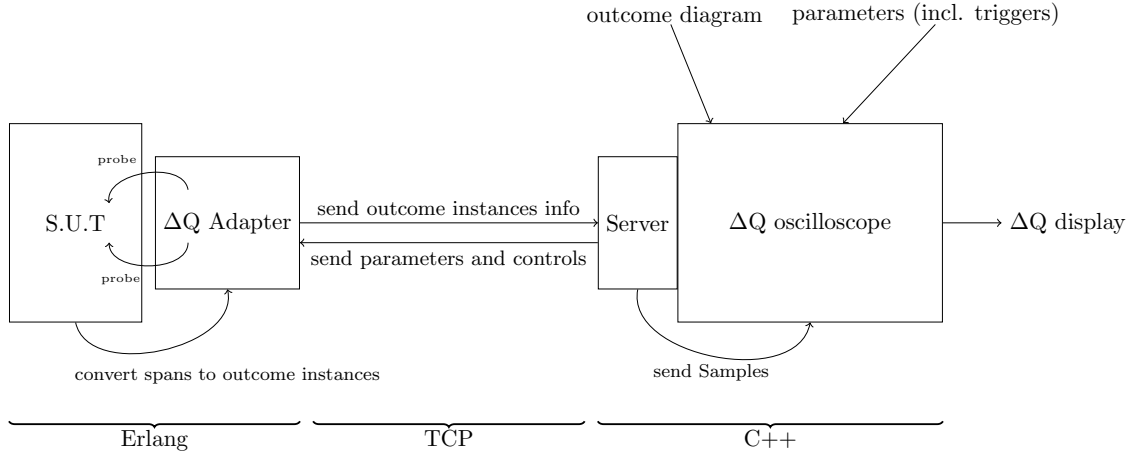


Figure 3.4: Global system design diagram.

3.7 Triggers

The concept of triggers is key to the oscilloscope, much like an oscilloscope that has a trigger mechanism to capture periodic signals or investigate a transient event [17], the ΔQ oscilloscope has a similar mechanism that can recognize when an observed ΔQ violates certain conditions regarding required behaviour and record snapshots of the system.

Each time an observed ΔQ is calculated by the oscilloscope, it is checked against the requirements set by the user. If these requirements are not met, a trigger is fired and a snapshot of the system is saved to be shown to the user.

3.7.1 Snapshot

A snapshot of the system gives insights into the system before and after a trigger was fired. It gives the user a still of the system, as if it was frozen in time.


When an observed ΔQ is calculated for each component being observed in the oscilloscope, they are stored away. In case the probe is observing multiple components, the ΔQ which is the result of the outcome expressions is also stored.

If no trigger is fired, older observed (and calculated) ΔQ s are removed. In the case that a trigger is fired, the oscilloscope keeps recording ΔQ s without removing older ones, to allow the user to look at the state of the system before the trigger and after.

3.8 Polling window

To calculate a ΔQ , we take all the outcome instances that ended within a window of time from t_l to t_u , a lower and upper time bound.

Suppose we are at time t , the window we will display is the **window of time** $(t-1)_l - (t-1)_u$ with $t-1$ equal to $t-x$, and x the polling rate. This is to account for various overheads that need to be taken into consideration, which could be network overhead, the wrapper overhead, C++ latency ... Imagine multiple outcome instances that are ended at a time slightly lower but close to t , and due to the overheads the messages arrives at a time slightly higher but close to t , the outcome instance would not be taken into consideration for the calculation of a ΔQ .



The polling window then advances every x seconds setting the new window:

$$\text{From: } (t-1)_l, (t-1)_u \xrightarrow{t+1} t_l, t_u.$$

$$\text{Where: } t_l = (t-1)_u \text{ and } t_u = (t-1)_u + x$$

Chapter 4

Oscilloscope: User level concepts

The following chapter gives insights on the user level concepts of ΔQSD in the oscilloscope. They are the concepts needed by the user to understand how the oscilloscope works.

- We first provide insights into how ΔQSD was implemented in the oscilloscope, the parameters that define a probe's ΔQ , its representation and what can be done with ΔQ s. We show how probe's $\Delta Q(s)$ will be shown in the oscilloscope.
- We then provide a language to write outcome diagrams based on an already existing syntax.
- Lastly, we explain the different controls present on the oscilloscope dashboard.

4.1 ΔQSD implementation

Originally, $\Delta Q(x)$ denotes the probability that an outcome occurs in a time $t \leq x$, defining then the "intangible mass" of such IRV as $1 - \lim_{x \rightarrow \infty} \Delta Q(x)$. We then extend the original definition to fit real time constraints, needing to calculate ΔQ s continuously.

For a given probe, $\Delta Q(t_l, t_u, dMax)$ is the probability that the time of series with n outcome instances with end time $t_l \leq t_e \leq t_u$, an outcome or probe occurs in time $t \leq dMax$.

4.1.1 Internal representation of a ΔQ

We provide a ΔQ class to calculate the ΔQ of a probe between a lower time bound t_l and an upper time bound t_u .

The ΔQ can be calculated in various ways:

Observed ΔQ The first way is by having n collected outcome instances between t_l and t_u , calculating its PDF and then calculating the *empirical cumulative distribution function* (ECDF) based on its PDF. This is called the **Observed ΔQ** .

Calculated ΔQ A ΔQ can also be calculated by performing operations which are the result of outcome expressions on two or more ΔQ s, the notion of outcome instances is then lost between calculations, as the interest shifts towards calculating the resulting PDFs and ECDFs. This is called the **Calculated ΔQ** .

4.1.2 dMax

The key concept of ΔQSD is having a maximum delay after which we consider that the execution is failed, this is represented in a prove as $dMax$. The user defines, for each prove the maximum delay its execution can have.

Setting a maximum delay for an prove is not a job that can be done one-off and blindly, it is something that is done with an underlying knowledge of the system inner-workings and must be thoroughly fine tuned during the execution of the system by observing the resulting distributions of the obtained ΔQ s.

We define in our oscilloscope a formula to dynamically define a maximum delay:

$$dMax = \Delta_T * N \quad (4.1)$$

Where Δ_T is the bin width of the ΔQ PDF and ECDF and N their number of bins.

The user must choose both via a slider. N is in the range $[1, 1000]$. This is a good enough bound to allow for finer grained representation, or less precision if needed.

Some tradeoffs must though be acknowledged when setting these parameters, a higher number of bins corresponds to a higher number of calculations and space complexity, a lower $dMax$ may correspond to more failures. These are all tradeoffs that must be considered by the system engineer and set accordingly.




Figure 4.1: $\Delta Q: dMax = 50ms$, the CDF will stay constant when delay $> dMax$

dMax limitation

$dMax$ can **not** be lower than 0 milliseconds and will be rounded to the **nearest** integer, this is a limitation of Erlang `send_after` function which only accepts integers and milliseconds values.

4.1.3 QTA

A simplified QTA is defined for probes. We define 4 points for the step function at 25, 50, 75 percentiles and the maximum amount of failures accepted for an observable. An observed ΔQ will calculate that based on the samples collected.

4.1.4 Confidence bounds

To observe the stationarity of a system we must observe a window of ΔQ s of an observable and calculate confidence bounds over said windows. The bounds can be updated dynamically by inserting or removing a ΔQ , this allows us to consider a small window of execution rather than observing the whole execution.




Figure 4.2: Upper and lower bounds (dashed, red) of the mean (blue) of multiple ΔQ s. In a system that behaves linearly, the bounds will be close to the mean, once the overload is approaching, or a system is showing behaviour that diverges from a linear one, the bounds will be larger.

4.2 ΔQ display

An probe's displayed graph must contain the following functions:

- The mean and confidence bounds of a window of previous ΔQ s.
- The observed $\Delta Q(t_l, t_u, dMax)$.
- If applicable, the calculated ΔQ from the components showing the causal links of a probe.
- Its QTA (if defined).

This allows for the user to observe if a ΔQ has deviated from normal execution, analyse its stationarity, nonlinearity and observe its execution.

4.3 Outcome diagram

An abstract syntax for outcome expressions and consequently outcome diagrams has already been defined in a previous paper [3], nevertheless, the oscilloscope provides additional features not included in the original syntax and, moreover, needs a textual way to define an outcome diagram.

We define thus a grammar to create an outcome diagram in our oscilloscope, our grammar is a textual interpretation of the abstract syntax.

4.3.1 Probes

To attach probes in the oscilloscope, the user must define outcomes and probes that observe outcome expressions.

Outcome

In our system an outcome is defined with its name

```
... = outcomeName;
```

Outcomes cannot be defined on their own and must be observed by other probes, we explain below how.

Probes containing outcome expressions

A probe can contain one outcome or a sequence of causally linked components. The user can define as many probes that contain outcome expressions as they want, they have to be declared as follows:

```
probe = component [-> component2];  
probe2 = newComponent -> anotherComponent;
```

These probes can be reused in other probes or in the system by adding "s:" (subsystem) before they are used.

```
probe3 = s:probe -> s:probe2;
```

The lines defining these probes must be semicolon terminated.

4.3.2 Operators, outcome expressions

To build a system, we must define the relations between outcomes and outcome expressions, below is how they can be defined.

First-to-finish, all-to-finish and probabilistic choice must contain at least two components, this is because the operations to calculate the *calculated* ΔQ rely on using the CDF of the components that define the operator.

Causal link

A causal link between two components can be defined by a right arrow from `component_i` to `component_j`

```
component_i -> component_j
```

All-to-finish operator

An all-to-finish operator needs to be defined as follows:

```
a:name(component1, component2...)
```

First-to-finish operator

A first-to-finish operator needs to be defined as follows.

```
f:name(component1, component2...)
```

Probabilistic choice operator

A probabilistic choice operator needs to be defined as follows:

```
p:name[probability_1, probability_2, ... probability_i](component_1,  
  ↪ component_2, ..., component_i)
```

In addition to being comma separated, the number of probabilities inside the brackets must match the number of components inside the parentheses. For n probabilities p_i , $0 < p_i < 1$, $\sum_{i=0}^n p_i = 1$

4.3.3 Limitations

Our system has a few limitations compared to the theoretical applications of ΔQ , namely, no cycles are allowed in the definition of a system.

```
probe = s:probe_2;  
probe_2 = s:probe;
```

The above example is not allowed and will raise an error when defined.

4.4 Dashboard

The dashboard is devised of multiple sections where the user can interact with the oscilloscope, create the system, observe the behaviour of its components, set triggers.

4.4.1 Sidebar

The sidebar has multiple tabs, we explain here the responsibility of each one.

System/Handle plots tab

System creation In this tab the user can create its system using the grammar defined before, he can save the text he used to define the system or load it, the system is saved to a file with any extension, we nevertheless define an extension to save the system to, the extension `.dq`. If the definition of the input is wrong, he will be warned with a pop up giving the error the parser generator encountered in the creation of a system.

Adding a plot Once the system is defined, the user can choose the probes he wants to plot. They can select multiple probes per plot and display multiple plots on the oscilloscope window.

Polling rate The user can choose the polling rate of the system: How often ΔQ s are calculated and displayed in the oscilloscope.

Editing a plot By clicking onto a plot that is being shown, the user can choose to add or remove probes to and from it. Multiple probes can be selected to either be removed or added.

Parameters tab

In this tab, the user can define parameters for the probes they have defined.

Set a QTA The user is given the choice to set a QTA for a given observable, they have 4 fields where they can fill in which correspond to the percentiles and the maximum amount of failures allowed, they can change this dynamically during execution.

dMax, bins The user has a slider which goes from -10 to 10, where they can set the parameters we explained previously, n , the exponent of $\Delta_{tbase} \cdot 2^n$ and the bins N . When these informations are saved by the user, the new $dMax$ is transmitted to the stub and saved for the selected observable.

Triggers tab

In the triggers tab the user can set triggers and observe the snapshots of the system.

Set triggers The user can set which triggers to fire for the probes they desire, they are given checkboxes to decide which ones to set as active or not (by default, the triggers are deactivated).

Fired triggers Once a trigger is fired, the system start a timer, during which all probes start recording the observed ΔQ s (and the calculated ones if applicable) without discarding older ones. Once the timer expires, the snapshot is saved for the user in the triggers tab. In the dashboard, it indicates when the trigger was fired (timestamp) and the name of the probe which fired it.

4.4.2 Plots window

To the left, the main window shows the plots of the probes being updated in real time.

4.4.3 Stub controls

Below the sidebar, two buttons are present, these buttons communicate to the stub.

The **start stub** button sends a message to the stub, telling it to start sending spans.
The **stop stub** button stops it.

4.5 Triggers

There are two available triggers which can be selected by the user, the triggers are evaluated on the **observed ΔQ** .

4.5.1 Load

A trigger on an observed ΔQ can be fired if the amount of outcome instances received in a polling rate is greater than what the user defines:

$$\#instances(\Delta Q(t_l, t_u, dMax)) > \text{maxAllowedInstances}$$

4.5.2 QTA

A trigger on an observed ΔQ can be fired if:

$$\Delta Q_{obs} > \text{observableQTA}$$

Chapter 5

Oscilloscope: implementation

The following chapter gives a more technical description of the oscilloscope.

- We provide a more in-depth look at the Δ QSD concepts introduced in the previous chapter.
- We then explain how the wrapper works, its API and the underlying mechanism that let us export outcome instances to the oscilloscope.
- Next we give a technical explanation of the parser generator we used to parse the outcome diagram syntax.
- Lastly, we briefly talk about the dashboard graphical framework.

5.1 Δ QSD implementation

A probe's Δ Q can be represented internally by a PDF and displayed as an ECDF. Here is how both can be calculated given n outcome instances.

PDF

We approximate the PDF of the a Δ Q via a histogram. We partition the values into N bins of equal width, this is required to ease future calculations. Given $[x_i, x_{i+1}]$ the interval of a bin i , where $x_i = i\Delta x$, and $\hat{p}(x_i)$ the value of the PDF at bin i , for n bins:

$$\begin{cases} \hat{p}(i) = \frac{n_i}{n}, \text{ if } i \leq n \\ \hat{p}(i) = 0, \text{ if } i > n \end{cases} \quad (5.1)$$

Where n_i the number of successful outcome instances whose elapsed time is contained in the bin i , n the total number of instances.

ECDF

The value $\hat{f}(x_i)$ of the ECDF at bin i with n bins can be calculated as:

$$\begin{cases} \hat{f}(i) = \sum_{j=1}^i \hat{p}(j), & \text{if } i \leq n \\ \hat{f}(i) = \hat{f}(x_n), & \text{if } i > n \end{cases} \quad (5.2)$$

5.1.1 dMax

We introduced $dMax$ in the previous section, we provide here the full equation that allows $dMax$ to be calculated:

$$dMax = \Delta_{base} * 2^n * N \quad (5.3)$$

Where:

- Δ_{base} represents the base width of a bin, equal to 1ms.
- N the number of bins.

5.1.2 Operations

In a previous section we talked about the possible operations that can be performed on and between ΔQ s, the time complexity of FTF, ATF and PC is trivially $\mathcal{O}(N)$ where N is the number of bins. As to convolution, the naïve way of calculating convolution has a time complexity of $\mathcal{O}(N^2)$, this quickly becomes a problem as soon as the user wants to have a more fine-grained understanding of a component. Below we present two ways to perform convolution.

Convolution

Naïve convolution Given two ΔQ binned PDFs f and g , the result of the convolution $f \otimes g$ is given by [18]:

$$(f \otimes g)[n] = \sum_{m=0}^N = f[m]g[n - m] \quad (5.4)$$

Fast Fourier Transform Convolution FFTW (Fastest Fourier Transform in the West) is a C subroutine library [19] for computing the discrete Fourier Transform in one or more dimensions, of arbitrary input size, and of both real and complex data. We use FFTW in our program to compute the convolution of ΔQ s. We adapt our script from an already existing one found on GitHub. [20]

Whilst the previous algorithm is far too slow to handle a high number of bins, convolution leveraging Fast Fourier Transform (FFT) allows us to reduce the amount of calculations to $\mathcal{O}(n \log n)$. This is why the naïve convolution algorithm is not used. We will analyse the time gains in a later chapter.

FFT and naïve convolution produce the same results in our program barring ε differences (around 10^{-18}) in bins whose result should be 0.

FFTs algorithms are plenty, the choice of the one to use is left up to the subroutine via the parameter `FFTW_ESTIMATE` [21].

Arithmetical operations

We can apply a set of arithmetical operations between ΔQ s ECDFs, and on a ΔQ .

Scaling (multiplication) A ΔQ can be scaled w.r.t a constant $0 \leq j \leq 1$. It is equal to binwise multiplication on ECDF bins.

$$\hat{f}_r(i) = \hat{f}(i) \cdot j \quad (5.5)$$

Operations between ΔQ s Addition, subtraction and multiplication can be done between two ΔQ of equal bin width (but not forcibly of equal length) by calculating the operation between the two ECDFs of the ΔQ s:

$$\Delta Q_{AB}(i) = \hat{f}_A(i)[\cdot, +, -]\hat{f}_B(i) \quad (5.6)$$

5.1.3 Confidence bounds

To observe the stationarity of a system we must observe a window of ΔQ s of an observable and calculate confidence bounds over said windows. We present here the formulae required to give such bounds with 68% confidence level.

For a bin i and an ECDF j , the mean of the bin over a window is:

$$\mu_i = \frac{1}{n_i} \sum_{j=1}^{n_i} x_{ij} \quad (5.7)$$

Where x_{ij} is a bin's i value for an ECDF j . Its variance:

$$\sigma_i^2 = \frac{1}{n_i} \sum_{j=1}^{n_i} x_{ij}^2 - \mu_i^2 \quad (5.8)$$

The confidence intervals CI_i for a bin i can then be calculated as:

$$CI_i = \mu_i \cdot \frac{\sigma_i}{\sqrt{n_i}} \quad (5.9)$$

5.1.4 Rebinning

Rebinning refers to the aggregation of multiple bins of a bin width i to another bin width j . Operations between ΔQ s can be done on ΔQ s that have the same bin width, this is why it is fundamental that all probes have a common Δ_{tbase} . This allows for fast rebinning to a common bin width.


Given two ΔQ s ΔQ_i , ΔQ_j :

$$\Delta_{Tij} = \max \{ \Delta_{Ti}, \Delta_{Tj} \}$$

and the PDF of the rebinned ΔQ at bin b , from the original PDF of n bins, where $k = \frac{\Delta T_j^i}{\Delta T_j}$.

$$p'_b = \sum_{n=b \cdot k}^{b+1 \cdot k - 1} p_n, \quad b = 0, 1, \dots, \lceil \frac{N}{k} \rceil \quad (5.10)$$

We perform rebinning to a higher bin width for a simple reason, while this leads to loss of information for the bin with the lowest bin width, rebinning to a lower bin width would imply inventing new values for the ΔQ with the highest bin width.



5.2 Wrapper

The wrapper, called `dqsd_otel` is a rebar3 application built to replace OpenTelemetry calls and create custom spans, it is designed to be paired with the oscilloscope to observe an erlang application.

5.2.1 API

The wrapper functions to be used by the user are made to replace OpenTelemetry calls to macros as for `?start_span` and `?with_span` and `?end_span`. This is to make the wrapper less of an encumbrance for the user.

Moreover, the wrapper will always start OpenTelemetry spans but only start custom spans if the stub has been activated. The wrapper can be activated by the oscilloscope by pressing the "start wrapper" button and can be stopped via the "stop wrapper" button.

start_span/1, start_span/2

```
start_span/1: -spec start_span(binary()) -> {opentelemetry:span_ctx(),
  ↳ pid() | ignore}.
start_span/2: -spec start_span(binary(), map()) ->
  ↳ {opentelemetry:span_ctx(), pid() | ignore}.
```

Parameters:

- Name: Binary name of the probe.

- Attributes: The OpenTelemetry span attributes (Only for `start_span/2`).

`start_span` incorporates OpenTelemetry `?start_span(Name)` macro.

Return: The function returns either:

- `{SpanCtx, span_process_PID}` if the wrapper is active and the observable's *dMax* has been set.
- `{SpanCtx, ignore}` if one of the two previous conditions was not respected.

With `SpanCtx` being the context of the span created by OpenTelemetry.

`with_span/1, with_span/2`

```
with_span/1: -spec with_span(binary(), fun(() -> any())) -> any().
with_span/2: -spec with_span(binary(), map(), fun(() -> any())) ->
  ↳ any().
```

Parameters:

- Name: Binary name of the probe.
- Fun: Zero-arity function representing the code of block that should run inside the `?with_span` macro.
- Attributes: The OpenTelemetry span attributes (Only for `with_span/3`).

`with_span` incorporates OpenTelemetry `with_span` macro.

Return: `with_span` returns what `Fun` returns (`any()`).

`end_span`

```
-spec end_span(opentelemetry:span_ctx(), pid() | ignore) -> ok |
  ↳ term().
```

Parameters:

- `SpanCtx`: The context of the span returned by `start_span`.
- `Pid`: `span_process_PID` || `ignore`.

As is the case for `start_span`, `end_span` incorporates an OpenTelemetry macro, in this case `?end_span(Ctx)`.

`fail_span`

```
-spec fail_span(pid() | ignore) -> ok | term().
```

Parameter:

- Pid: `ignore || span_process_PID`.

`fail_span` does not incorporate any OpenTelemetry macro, it is let up to the user to decide how to handle failures in execution.

`span__process`

`span_process` is the process, spawned by `start_span`, responsible for handling the `end_span`, `fail_span`, `timeout` messages.

Upon being spawned, the process starts a timer with time equal to the `dMax` set by an user for the observable being observed, thanks to `erlang:send_after`. when the timer runs out, it sends a `timeout` message to the process.

The process can receive three kinds of messages:

- `{end_span, end_time}`: This will send a custom span to the oscilloscope with the start and end time of the execution of the observable.
- `{fail_span, end_time}`: This will send a custom span to the oscilloscope indicating that an execution of an observable has failed.
- `{timeout, end_time(StartTime + dMax)}`: If the program hasn't ended the span before `dMax`, the timer will send a `timeout` message and it will send a custom span to the oscilloscope indicating that an execution of an observable has taken $> dMax$.

The process is able to receive one and only message, if the execution times out and subsequently the span is ended, the oscilloscope will not be notified as the process is defunct. This is assured by Erlang documentation:

If the message signal was sent using a process alias that is no longer active, the message signal will be dropped. [22]

5.2.2 Handling outcome instances

To create outcome instances of a probe we must obtain three important informations:

- Its name.
- The time when the span was started.
- Its `dMax`.

Their start time and end time are supplied by calling this function:

```
StartTime/EndTime = erlang:system_time(nanosecond).
```

The name is given when starting a span and the `dMax` is stored in a dictionary in the wrapper.

The outcome instance is created only if two conditions are met: the wrapper has been set as active and the user set a timeout for the probe, the functions will spawn a `span_process` process, passing along all the necessary informations.

Once the span is subsequently ended/timed out/failed, the function `send_span` creates a message carrying all the informations and sends it to the C++ server. The formatting of the messages is the following:

```
n:Observed name, b: Start time (beginning), e: End time (end or  
↪ deadline), s: The status
```

5.2.3 TCP connection

The wrapper is composed of two `gen_server` which handle communication to and from the oscilloscope. This `gen_server` behaviour allows the wrapper to send spans asynchronously to the oscilloscope.

TCP server

The TCP server listens by default on localhost at port 8081, the user can define a port and ip to listen at.

The oscilloscope can send commands to the stub, these commands are:

- `start_stub`: This command sets the stub as active, it can now send spans to the oscilloscope if the items are defined.
- `stop_stub`: This commands sets the stub as inactive, it will no longer send spans to the oscilloscope.
- `set_timeout;probeName;timeout`: This command indicates to the stub to set the *dMax* for a probe to timeOut, a limit of the stub is that erlang:send_after does not accept floats as timeouts, so the timeout will be rounded to the nearest integer.

TCP client

The TCP client allows the stub to send the spans to the oscilloscope, by default the oscilloscope is on localhost:8080, but that can be changed by the user.

The client connects over TCP to the oscilloscope and opens a socket where it can send the spans.

5.3 Parser

To parse the system, we use the C++ ANTLR4 (ANother Tool for Language Recognition) library.

5.3.1 ANTLR

ANTLR is a parser generator for reading, processing, executing or translating structured text files. ANTLR generates a parser that can build and walk parse trees [23].

ANTLR is just one of the many parsers generators available in C++ (flex/bison, lex, yacc), although it presents certain limitations, its generated code is simpler to handle and less convoluted with respect to the other possibilities.

ANTLR uses Adaptive LL(*) (*ALL*(*)) parser, namely, it will move grammar analysis to parse-time, without the use of static grammar analysis. [24]

5.3.2 Grammar

ANTLR provides a yacc-like metalanguage to write grammars. Below, is the grammar for our system: [cite]

```
grammar DQGrammar;

PROBE_ID: 's';
BEHAVIOR_TYPE: 'f' | 'a' | 'p';
NUMBER: [0-9]+('.'[0-9]+)?;
IDENTIFIER: [a-zA-Z_] [a-zA-Z0-9_]*;
WS: [ \t\r\n]+ -> skip;

// Parser Rules
start: definition* system? EOF;

definition: IDENTIFIER '=' component_chain ' ';
system: 'system' '=' component_chain ' ';

component_chain
    : component ('->' component)*
    ;

component
    : behaviorComponent
    | probeComponent
    | outcome
    ;

behaviorComponent
    : BEHAVIOR_TYPE ':' IDENTIFIER ('[' probability_list ']')? '('
      ↪ component_list ')'
    ;

probeComponent
    : PROBE_ID ':' IDENTIFIER
    ;
```

```
probability_list: NUMBER (',' NUMBER)+;
component_list: component_chain (',' component_chain)+;

outcome: IDENTIFIER;
```

Limitations

A previous version was implemented in Lark[cite], a python parsing toolkit. The python version was quickly discarded due to a more complicated integration between Python and C++. Lark provided Earley(SPPF) strategy which allowed for ambiguities to be resolved, which is not possible in ANTLR.

For example the following system definition presents a few errors:

```
probe = s -> a -> f -> p;
```

While Lark could correctly guess that everything inside was an outcome, ANTLR expects ":" after "s, a, f" and "p", thus, one can not name an outcome by these characters, as the parsers generator thinks that an operator or a probe will be next.

5.4 Oscilloscope GUI

Our oscilloscope graphical interface has been built using the QT framework for C++. Qt is a cross-platform application development framework for creating graphical user interfaces. We chose Qt as we believe that it is the most documented and practical library for GUI development in C++, using Qt allows us to create usable interfaces quickly, while being able to easily pair the backend code of C++ to the frontend.

The interface is composed of a main window, where widgets can be attached to it easily. Everything that can be seen is customisable widgets. This allows for easy reusability, modification and removal without great refactoring due in other parts of the system.

In the photo below we can see each top level widget (a QWidget that contains other widgets) in the main window, the widgets could easily be switched to other places of the window or rearranged.




Figure 5.2: The oscilloscope displaying probes, the boxes represent a top level widget, which may contain other widgets inside.

Chapter 6

Application on synthetic programs

This section aims to provide an example of how the oscilloscope could be used to instrument an application, in this case, a synthetic one. We explain how the Δ QSD paradigm can be applied to explore tradeoffs in design and to gain more insights into a running system.

6.1 Sequential composition with M/M/1/K queues

Why M/M/1/K queues? An average component in a distributed system can be modeled as an M/M/1/K, due to the exponential inter-arrival rate of messages, the exponential distribution of the execution delay and the message buffer size of a component.

6.1.1 System composition

The system has two components `worker_1`, `worker_2`, the components are made of a buffer queue of size K and a worker process.

The system sends n messages per second following a Poisson distribution to `worker_1`'s queue, the queue then reduces its available buffer size.

The buffer notifies its worker, which then does N loops, which are defined upon start, of fictional work. The worker then passes a message to `worker_2`'s queue, which has another queue of same size, who passes the message to `worker_2`'s worker, which does the same amount of loops. When a worker completes its work, it notifies the queue, freeing one "message" from its buffer size.

If the queue's buffer is overloaded, it will drop the incoming message and consider the execution a failure.

A probe p is defined, which observes the execution from when the first message up until `worker_2` is done.





Figure 6.1: Outcome diagram of the M/M/1/K queue with the colored lines representing the probes that were inserted.

6.1.2 Determining parameters dynamically

We stated previously that determining parameters is something that must be done with an underlying knowledge of the system. The oscilloscope can provide knowledge of the system, here is an example of worker_1 and worker_2 as observed in the oscilloscope.

Imagine the engineer supposes the workers executions should take around 6.5 ms to complete, but doesn't actually know how long the executions should take. The engineer, after having set the required parameters observes in the following graph in the oscilloscope Section 6.1.2.

The oscilloscope shows the engineer that their assumptions do not correspond to the actual system ΔQ , the user can then modify the parameters to observe the actual system's behaviour. By setting $dmax$ to 25 ms, he can observe the worker's ΔQ s approaching 1.




(a) worker_1 and worker_2 ΔQ s plot with 6.5 ms $dMax$. (b) worker_1 and worker_2 ΔQ s plot with 25 ms $dMax$.

On the other hand, the engineer's assumption could have been what he truly expected


from the system, in this case, the oscilloscope tells him that the system is not behaving as expected.

Low Load At low load, we can observe in the oscilloscope how worker_1 and worker_2 mean ΔQ s overlap. This is expected, even under overload or dependency conditions, worker_1 and worker_2 should have the same ΔQ s.

If the system is not showing dependent behaviour, the probe **observed** ΔQ and **calculated** ΔQ should overlap. We can observe that in the graph below, as the mean CDF of both ΔQ s overlap.




(a) worker_1 (blue) and worker_2 (purple) ΔQ plots as shown in the oscilloscope



(b) probe ΔQ plots as shown in the oscilloscope

Early signs of overload Once the system is approaching overload, we can see the two means starting to diverge. As shown in a previous chapter, the observed (1) ΔQ of the probe is below the calculated (2) ΔQ . (explain better)



6.1.3 Triggers - QTA

Triggers and QTAs are an important part of the oscilloscope and ΔQSD , we will show here how they can be used and what ΔQ can tell about the state of the system.

Satisfying timeliness, a tale of typical queue behaviour ΔQ can be used to feel overload approaching and analyse a system's behaviour. To do so, we set a QTA for the probe. 25% of instances should take < 6.5 ms, 50% $< 12ms$, 75% $< 17ms$. We run the program and wait for triggers to be fired. To do so, we synthetically overload the CPU to degrade the system's performance.

From the ΔQ s we have in the snapshot we can observe:

1. **Business as usual:** The system is not showing signs of stress, the ΔQ is better than or between than the confidence bounds.
2. **Overload approaching:** The component is slowing down and the ΔQ is slightly worse than previously, but we can still feel overload is approaching.
3. **QTA violation:** The queue, overloaded, quickly violates the QTA.
4. **Get in trouble quickly, get out of it slowly:** The buffering filling up quickly and emptying slowly is a typical queue pattern, this can be observed on the oscilloscope

Figure 6.4: The snapshots fired during execution, we will select the snapshot triggered at 17:52:33

Triggered Snapshots:
Snapshot at: Fri May 23 17:52:07 2025 from probe
Snapshot at: Fri May 23 17:52:20 2025 from probe
Snapshot at: Fri May 23 17:52:26 2025 from probe
Snapshot at: Fri May 23 17:52:33 2025 from probe

For the following graphs, the **blue** line, annotated as **(1)** is the **observed ΔQ** , the **red** line, annotated as **(2)**, is the **calculated ΔQ** .

Business as usual

During normal execution, we can observe the system behaving correctly for a long time before the trigger is fired. This is what the user who connects the oscilloscope to a system which correctly behaves should see most of the time.




Figure 6.5: Normal execution of the system over 2 seconds before overload.

Overload approaching

We can remark the system approaching overload just before the system violates the QTA. While no triggers are yet fired, as the system is still respecting QTA, we can feel something is starting to go wrong.




Figure 6.6: The system is slowly degrading just seconds before overload.

QTA violation

Just moments after the system performance degrading, the observed ΔQ violates the QTA. We point in blue where the violations happen.

We can observe how the system ΔQ is dangerously close to the QTA bounds right after violating the QTA. This is because the queue will fill up quickly and empty slowly. We can observe in the next section how the repercussions of such violations linger on for some time after the violation happens.




Figure 6.7: Left: The system violates the QTA. Right: The system is slowly recuperating from overload.

Get in trouble quickly, get out of it slowly

By recording a snapshot the system **after** the violation happens, we can observe how long it takes to get out.





Figure 6.8: The system is slowly degrading just seconds before overload.

The system gets out of dangerous grounds **2 seconds** after violating the QTA, and still shows the consequences of overload 3 seconds after.



6.2 First to finish application

Next, we provide a synthetic application modeling an application that can be modeled by a first to finish operator

Why first to finish? Recall the previous FTF graph Figure 2.5. Assume a send request to "the cloud" that waits for a response or a timeout, it is modeled by a FTF operator. A sample resulting graph of "the cloud" is the one we showed previously.

[INSRET PGRAPH]

6.2.1 Using the wrong operator

What happens if the wrong operator is chosen to represent the causal relationships between the outcomes? What if the user believes that the system diagram is the one we presented before Figure 6.1? The result on the oscilloscope will clearly show that something is wrong!




Figure 6.10: *(Left)* FTF plot **with wrong outcome diagram definition** as shown in the oscilloscope. (1) Observed ΔQ . (2) Calculated ΔQ . *(Right)* FTF plot **with correct outcome diagram definition** as shown in the oscilloscope. Observed ΔQ and calculated ΔQ overlapping.

On the left, we can observe how the **calculated ΔQ** (2) is clearly greater than the **observed ΔQ** (1). A difference this drastic tells us that the proposed outcome diagram does not correctly represent the actual system. On the right, if no dependencies are present and the correct operator is chosen, the two graphs will overlap.

6.2.2 Introducing a slower component

Let us introduce a slower worker into the system, we introduce an artificial delay into worker_2 (about 20ms). If the oscilloscope works correctly, the paradigm operations are sound and no dependencies are present in the system, we should not see any difference in the observed and calculated ΔQ s of the FTF operator.




Figure 6.11: *(Left)* FTF plot of worker_1 and worker_2, observed and calculated ΔQ overlapping.

(Right) worker_1 (1) and worker_2 (2) ΔQ s.

The FTF plot correctly displays how worker_2 does not have an effect on the ftf plot.

Chapter 7

Performance study

This chapter evaluates the components and operations we introduced in previous sections, analysing their performances

7.1 Convolution performance

We implemented two versions of the convolution algorithm as described before, the naïve version and the FFT version. We compared their performance when performing convolution on two ΔQ s of equal bins.




Figure 7.1: Performance comparison of two convolution algorithms

As expected, the naïve version has a time complexity of $\mathcal{O}(n^2)$ and quickly scales with the number of bins, this is clearly inefficient, as a more precise ΔQ will result in a much slower program.


As for the FFT algorithm, it is slightly slower when the number of bins is lower than 100. This is due to the FFTW3 routine having slightly higher overhead.

7.2 Stub performance

We evaluated the performance of the stub to measure its impact in a normal execution, namely we tested the following calls which represent a normal usage of the wrapper.

- `start_span` \rightarrow `end_span`.
- `with_span` with the following function: `fun()` \rightarrow `ok`.
- `start_span` \rightarrow `fail_span`.

We ran the simulation for 25000 subsequent iterations, these are the results.




The overhead is minimal, around 10 microseconds on average to start and end/fail a span. The same cannot be said about with span, the increased overhead is nevertheless due to a function needing to be called inside it for it to record a span.

7.3 GUI plotting performance

We evaluated the performance of the GUI plotting routine for an observed ΔQ . The routine first prepares the ΔQ , creating vectors `QPointF` (a Qt class representing a point for a `QtChart`), representing the x and y values of the ΔQ s CDF. The vectors are created for the lower bound, the upper bound the mean of the window of ΔQ s and the observed ΔQ .

Then, once the vectors are prepared, Qt replaces the old points with the new points for every element being plotted.



The result scales up to 2 ms for 1000 bins, since multiple plots have to be shown, if the probe contains outcome expressions inside, the calculated ΔQ will need to be plotted too, the time increase will be twofold. If a user decides to have finer grained representations and multiple plots at once, decreasing (decreasing or increasing?) the polling rate will avoid the plots frame skipping.

Chapter 8

Conclusions and future work

The following project is the beginning of the ΔQ oscilloscope, our initial goal was to create an application to observe running distributed applications, namely, Erlang ones. A prototype was successfully created thanks to the following feats:

- The graphical dashboard for the ΔQ oscilloscope, built in C++, which allows real time display of ΔQ s for the probes inserted in the system.
- Fast convolution algorithms to perform statistical analysis on probes.
- The creation of a textual syntax to create outcome diagrams.
- The `dqsd_otel` Erlang wrapper to connect an OpenTelemetry instrumented Erlang app to the oscilloscope.

The user has full control over the system and can update it dynamically to add or remove probes, this allows full control of what the user decides to include, allowing a finer grained outcome diagram or a more general view of the system.

The oscilloscope and the Erlang can communicate via TCP socket connections to exchange outcome instances and probe parameters,

We showed how it can be useful in detecting early signs of overload many crucial features are still missing from the dashboard and it could require less code modifications in the Erlang side.

8.1 Future improvements

We believe the oscilloscope and the Erlang application can be drastically improved, the size of the project and its intended goal is too big to be encompassed in a single master thesis. We list here some improvements which could be made to both the oscilloscope and the wrapper.

8.1.1 Oscilloscope improvement

- The oscilloscope could be turned into a **web app**, we feel that a C++ oscilloscope is a good prototype and proof of concept, but its usability would be greater in a browser context. It would be great as a plugin for already existing observability platforms like Grafana.
- A wider selection of **triggers**, as of writing this thesis, only the QTA trigger and load are available, this is a limitation due to time constraints. Nevertheless, triggers can be easily implemented in the available codebase.
- **Better communication between stub - server - oscilloscope.** The current way of sending outcome instances may be a limiting factor under high load, if hundred of thousands of spans were to be sent, the current way the server and oscilloscope are tied together may throttle communications. TCP socket connections could quickly become the chokepoint which makes the oscilloscope temporarily unusable.

Future improvements on the server side could implement epoll system server calls to make the server more efficient; **Detaching server from client**, as of right now, the oscilloscope and the server are tied together, using ZeroMQ to assure real time server-client communications could be an interesting solution to explore.

- **Improve real time graphs.** The class QtCharts does not perform correctly with high frequencies update. Moreover, since we are plotting multiple series (from a minimum 4 to a maximum of 9) per probe, which allows up to 1000 bins per probe, the performance quickly degrades with more probes being displayed. A better graphing class for Qt could definitely improve the experience.
- **Saving probe parameters:** As of writing this thesis, there is no way to save the parameters one may have set.
- **Deconvolution:** An important aspect of Δ QSD, which was not introduced in this paper is deconvolution. It is used to check for infeasability in system desing. Since convolution has already been implemented, this could be integrated using the FFTW3 library.
- **Exporting graphs:** The graphs can only be observed in the oscilloscope and have no way to be exported to other programs via standard formats.
- **Many more:** This oscilloscope is just a start, if we were to list everything we may want to add, it would take many pages. What we provide is a sufficient enough basis to provide possibilities to observe a running system and understand the power of Δ QSD in analysing its behaviour.

8.1.2 Wrapper improvements

- As suggested by Bryan Naegele, a member of the observability group of Erlang, the wrapper, instead of working on top of OpenTelemetry, could be directly included inside the context of a span by using the ctx library [25], which provides deadlines for contexts, propagating the value in `otel_ctx`, making it available

to the OpenTelemetry span processor. Leveraging `erlang:send_after` as we already do, we could create outcome instances with telemetry events to handle successful executions and timeouts. The span processor will then be responsible for creating outcome instances, without creating the need for custom functions in the wrapper, like we have now.

8.1.3 Real applications

A flaw of the oscilloscope and wrapper is that they have not been tested on real applications, while their usefulness has been proven on synthetic applications, the lack of real life applications is a weakness.

8.1.4 Licensing limitations

Lastly, a notable limitation is created by **Qt**, namely, QtCharts. The usage of Qt does not allow us to release our project under BSD/MIT licenses, but rather a GPLv3 one (we cannot release it under LGPL due to QtCharts). [26]

Bibliography

- [1] Peter Van Roy and Seyed Hossein Haeri. *The Δ QSD Paradigm: Designing Systems with Predictable Performance at High Load. Full-day tutorial*. 15th ACM/SPEC International Conference on Performance Engineering. 2024. URL: <https://webperso.info.ucl.ac.be/~pvr/ICPE-2024-deltaQSD-full.pdf>.
- [2] Seyed H. Haeri et al. “Mind Your Outcomes: The Δ QSD Paradigm for Quality-Centric Systems Development and Its Application to a Blockchain Case Study”. In: *Comput.* 11.3 (2022), p. 45. DOI: 10.3390/COMPUTERS11030045. URL: <https://doi.org/10.3390/computers11030045>.
- [3] Seyed Hossein Haeri et al. “Algebraic Reasoning About Timeliness”. In: *Proceedings 16th Interaction and Concurrency Experience, ICE 2023, Lisbon, Portugal, 19th June 2023*. Ed. by Clément Aubert et al. Vol. 383. EPTCS. 2023, pp. 35–54. DOI: 10.4204/EPTCS.383.3. URL: <https://doi.org/10.4204/EPTCS.383.3>.
- [4] Peter Van Roy. *LINFO2345 lessons on Δ QSD*. Accessed: (19/05/2025). UCLouvain, 2023. URL: <https://www.youtube.com/watch?v=tF7fbU9Gce8>.
- [5] Neil J. Davies and Peter W. Thompson. *Δ QSD workbench - GitHub*. Accessed: (19/05/2025). 2022. URL: <https://github.com/DeltaQ-SD/dqsd-workbench>.
- [6] Erlang programming language. *Erlang tracing*. Accessed: (19/05/2025). 2024. URL: <https://www.erlang.org/doc/apps/erts/tracing.html>.
- [7] OpenTelemetry. *OpenTelemetry in Erlang/Elixir*. Accessed: (19/05/2025). 2025. URL: <https://opentelemetry.io/docs/languages/erlang/>.
- [8] Peter Thompson and Rudy Hernandez. *Quality Attenuation Measurement Architecture and Requirements*. Tech. rep. MSU-CSE-06-2. Sept. 2020. URL: <https://www.broadband-forum.org/pdfs/tr-452.1-1-0-0.pdf>.
- [9] OpenTelemetry. *What is OpenTelemetry?* Accessed: (19/05/2025). 2025. URL: <https://opentelemetry.io/docs/what-is-opentelemetry/>.
- [10] OpenTelemetry. *OpenTelemetry - Traces*. Accessed: (19/05/2025). 2025. URL: <https://opentelemetry.io/docs/concepts/signals/traces/>.
- [11] The Jaeger Authors. *Jaeger*. Accessed: (19/05/2025). 2025. URL: <https://www.jaegertracing.io/>.
- [12] Dotan Horovits. *From Distributed Tracing to APM: Taking OpenTelemetry & Jaeger Up a Level*. Accessed: (19/05/2025). 2021. URL: <https://logz.io/blog/monitoring-microservices-opentelemetry-jaeger/>.
- [13] Sampath Siva Kumar Boddeti. *Tracing Made Easy: A Beginner’s Guide to Jaeger and Distributed Systems*. Accessed: (19/05/2025). 2024. URL: <https://openobserve.ai/blog/tracing-made-easy-a-beginners-guide-to-jaeger-and-distributed-systems/>.

- [14] OpenTelemetry. *Instrumentation for OpenTelemetry Erlang/Elixir*. Accessed: (19/05/2025). 2025. URL: <https://opentelemetry.io/docs/languages/erlang/instrumentation/>.
- [15] OpenTelemetry. *Active spans, C++ Instrumentation*. Accessed: (19/05/2025). 2025. URL: <https://opentelemetry.io/docs/languages/cpp/instrumentation/>.
- [16] Hazel Weakly. *OpenTelemetry Challenges: Handling Long-Running Spans*. Accessed: (21/05/2025). 2024. URL: <https://thenewstack.io/opentelemetry-challenges-handling-long-running-spans/>.
- [17] KeySight. *What is an Oscilloscope Trigger?* Accessed: (23/05/2025). 2022. URL: <https://www.keysight.com/used/id/en/knowledge/glossary/oscilloscopes/what-is-an-oscilloscope-trigger>.
- [18] Steven B. Damelin and Willard Miller Jr. *The Mathematics of Signal Processing*. USA: Cambridge University Press, 2012. ISBN: 1107601045.
- [19] FFTW3. *Fastest Fourier Transform in The West*. Accessed: (19/05/2025). 2025. URL: <https://www.fftw.org/>.
- [20] Jeremy Fix. *FFTConvolution*. Accessed: (21/05/2025). 2013. URL: https://github.com/jeremyfix/FFTConvolution/blob/master/Convolution/src/convolution_fftw.h.
- [21] *Planning-rigor flags*. Accessed: (23/05/2025). URL: <https://www.fftw.org/doc/Planner-Flags.html>.
- [22] Erlang System Documentation. *Signals/Sending signals*. Accessed: (24/05/2025). 2025. URL: https://www.erlang.org/doc/system/ref_man_processes.
- [23] ANTLR. *What is ANTLR4?* Accessed: (19/05/2025). 2025. URL: <https://www.antlr.org/>.
- [24] Terence Parr and Kathleen Fisher. “LL(*): the foundation of the ANTLR parser generator”. In: *Proceedings of the 32nd ACM SIGPLAN Conference on Programming Language Design and Implementation, PLDI 2011, San Jose, CA, USA, June 4-8, 2011*. Ed. by Mary W. Hall and David A. Padua. ACM, 2011, pp. 425–436. DOI: 10.1145/1993498.1993548. URL: <https://doi.org/10.1145/1993498.1993548>.
- [25] Tristan Sloughter. *ctx*. Accessed: (21/05/2025). 2023. URL: <https://github.com/tsloughter/ctx>.
- [26] The Qt Company. *Add-ons available under Commercial Licenses, or GNU General Public License v3*. Accessed: (23/05/2025). 2025. URL: <https://doc.qt.io/qt-5/qtmodules.html>.

UNIVERSITÉ CATHOLIQUE DE LOUVAIN
École polytechnique de Louvain

Rue Archimède, 1 bte L6.11.01, 1348 Louvain-la-Neuve, Belgique | www.uclouvain.be/epl